

CSE508 Information Retrieval

ASSIGNMENT-3 Report

Divyansh Mishra

2021042

In order to account for the computational power required for this assignment we used kaggle.
Reading the json files:

```
import pandas as pd
import gzip
import feather
import json

def parse(path):
    with open(path, 'r') as g:
        for l in g:
            yield json.loads(l)

def getDF(path):
    i = 0
    df = {}
    for d in parse(path):
        df[i] = d
        i += 1
    return pd.DataFrame.from_dict(df, orient='index')

df = getDF('/kaggle/input/amazon-rev-data/amazon-review-data/Electronics_5.json/Electronics_5.json')
```

We chose the product console for this assignment.

```
df_cat = df_meta[df_meta['title'].str.contains('console', case=False)]
df_cat.drop_duplicates(subset=['asin'], inplace=True)
print(len(df_cat))
# df_headphones.head
df_filtered_reviews = df[df['asin'].isin(df_cat['asin'])]
print(len(df_filtered_reviews))
# df_merged = df_filtered_reviews.merge(df_cat, on='asin', how='inner')
# df_filtered_reviews=df_merged
# df_filtered_reviews = pd.merge(df, df_cat, on='asin', how='inner')
# print(df_filtered_reviews.head)
```

```
Number of reviews: 5836
Average Rating Score: 4.3658327621658675
Number of Unique Products: 145
Number of Good Ratings: 5255
Number of Bad Ratings: 581
```

```
Number of Reviews corresponding to each Rating:
overall
1.0      368
2.0      213
3.0      344
4.0      902
5.0     4009
Name: count, dtype: int64
```

```
# Apply preprocessing functions
df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(remove_html_tags)
df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(remove_accented_chars)
df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(remove_special_characters)
df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(lemmatize_text)
df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(normalize_text)
# df_filtered_reviews['reviewText'] = df_filtered_reviews['reviewText'].apply(ch)
```

```
Count of Ratings for Most Positively Reviewed Console Over 5 Consecutive Years:
Year 2014: 749 ratings
Year 2015: 1144 ratings
Year 2016: 1296 ratings
Year 2017: 1201 ratings
Year 2018: 769 ratings
```



The EDS that we performed gave us these values: and the product ID of the most positively reviewed product on the basis of average overall for each asin.

Top 20 Most Reviewed Brands:

brand	
Plugable	606
ORIA	604
BenQ	576
Logitech	374
Panlong	271
Belkin	242
Cellnorth Electronics	237
by\n \n TOMSENN	186
Cisco	181
Asunflower	172
HDE	169
StarTech	139
Fosmon	129
UGREEN	129
Alienware	114
hossen	109
Creative	84
E-sds	84
Edimax	81
Pyle	80

Name: count, dtype: int64

Top 20 Least Reviewed Brands:


brand	
TRENDnet	6
IOGEAR	6
InterlinkAuckland	6
MOCREO	6
ULBRE	6
Bose	5
Stuff4	5
Hausbel	5
Rocketfish	5
2gig	5
Crosley	5
Like	5
Inland	5
Importer520	5
Zendure	5
Marantz	5
Hercules	5
Abrams	5
PowerMonkey	4
elexa consumer products	4

Name: count, dtype: int64

ASIN: B01EAYYJAS

	overall	vote	verified	reviewTime	reviewerID	asin	style	\
5799	5.0	NaN	True	07 15, 2018	A208EYXZZ1PPYL	B01EAYYJAS	NaN	
5800	5.0	NaN	True	06 23, 2018	A2VS5SQROVJK41	B01EAYYJAS	NaN	

Distribution:



```
1: # Group the reviews by 'reviewYear' and count the number of reviews for each year
reviews_per_year = df_filtered_reviews['reviewYear'].value_counts()

# Find the year with the maximum number of reviews
max_reviews_year = reviews_per_year.idxmax()
max_reviews_count = reviews_per_year.max()

# Report the year with the maximum reviews
print(f"The year with the maximum reviews is {max_reviews_year} with {max_reviews_count} reviews.")
```

The year with the maximum reviews is 2016 with 1296 reviews.

Year with the Highest Number of Customers (Considering Verified Reviews Only):
Year: 2016, Number of Customers: 1163

We the the TF-IDF for feature engineering

The ML models that we used were: "Logistic Regression", "Random Forest", "Support Vector Classifier", "Multinomial Naive Bayes", "Gradient Boosting Classifier"

```
Logistic Regression Classification Report:
              precision    recall  f1-score   support

   Average      0.00      0.00      0.00        99
     Bad       0.78      0.31      0.45       143
     Good       0.86      1.00      0.93      1217

 accuracy              0.86      1459
 macro avg              0.55      0.44      0.46      1459
 weighted avg              0.80      0.86      0.82      1459
```

=====

Training Random Forest...

```
Random Forest Classification Report:
              precision    recall  f1-score   support

   Average      0.29      0.02      0.04        99
     Bad       0.83      0.21      0.34       143
     Good       0.85      0.99      0.92      1217

 accuracy              0.85      1459
 macro avg              0.66      0.41      0.43      1459
 weighted avg              0.81      0.85      0.80      1459
```

=====

Training Support Vector Classifier...

```
Support Vector Classifier Classification Report:
              precision    recall  f1-score   support

   Average      0.50      0.01      0.02        99
     Bad       0.76      0.43      0.55       143
     Good       0.88      0.99      0.93      1217

 accuracy              0.87      1459
 macro avg              0.71      0.48      0.50      1459
 weighted avg              0.84      0.87      0.83      1459
```

=====

Training Multinomial Naive Bayes...

```
Multinomial Naive Bayes Classification Report:
              precision    recall  f1-score   support

   Average      0.00      0.00      0.00        99
     Bad       1.00      0.01      0.01       143
     Good       0.83      1.00      0.91      1217

 accuracy              0.83      1459
 macro avg              0.61      0.34      0.31      1459
 weighted avg              0.79      0.83      0.76      1459
```

=====

Training Gradient Boosting Classifier...

```
/opt/conda/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1
ed and being set to 0.0 in labels with no predicted samples. Use `zero_divis
_warn_prf(average, modifier, msg_start, len(result))
/opt/conda/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1
ed and being set to 0.0 in labels with no predicted samples. Use `zero_divis
_warn_prf(average, modifier, msg_start, len(result))
/opt/conda/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1
ed and being set to 0.0 in labels with no predicted samples. Use `zero_divis
_warn_prf(average, modifier, msg_start, len(result))
```

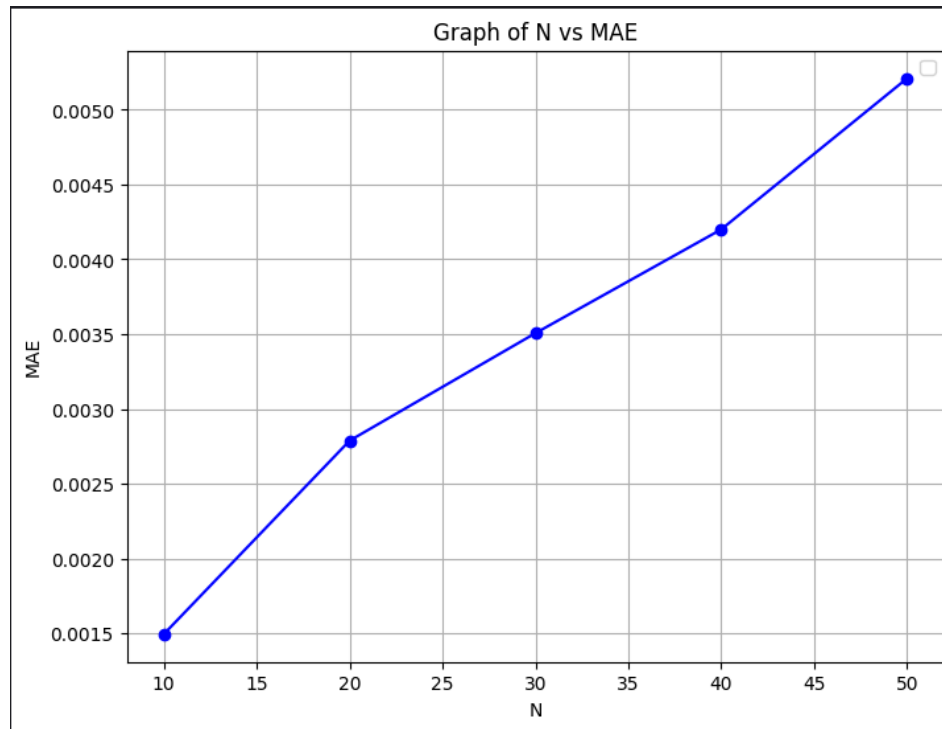
```
Gradient Boosting Classifier Classification Report:
              precision    recall  f1-score   support

   Average      0.67      0.02      0.04        99
     Bad       0.68      0.28      0.40       143
     Good       0.86      0.99      0.92      1217


 accuracy              0.86      1459
 macro avg              0.74      0.43      0.45      1459
 weighted avg              0.83      0.86      0.81      1459
```

=====

```
MAE AT N= 10 is 0.0014940755322140561
MAE AT N= 20 is 0.002787501995090178
MAE AT N= 30 is 0.0035063424866803978
MAE AT N= 40 is 0.004200063842885699
MAE AT N= 50 is 0.005209761881237032
```



```
MAE for Item-Item Recommender System with N=10: 0.043581357877376055
MAE for Item-Item Recommender System with N=20: 0.03534285149688765
MAE for Item-Item Recommender System with N=30: 0.045681991492682136
MAE for Item-Item Recommender System with N=40: 0.03744580955059169
MAE for Item-Item Recommender System with N=50: 0.034672660804268354
```



```
# Calculate the sum of ratings for each product across all users
product_sum_ratings = df_filtered_reviews.groupby('asin')['overall'].sum()

# Sort the products based on sum of ratings in descending order
top_10_products = product_sum_ratings.sort_values(ascending=False).head(10)

# Print the top 10 products by sum of ratings
print("Top 10 Products by User Sum Ratings:")
for rank, (product, rating_sum) in enumerate(top_10_products.items(), 1):
    print(f"{rank}. Product ID: {product}, Sum of Ratings: {rating_sum}")
```

```
Top 10 Products by User Sum Ratings:
1. Product ID: B01E16J6RQ, Sum of Ratings: 2762.0
2. Product ID: B00AQM8586, Sum of Ratings: 2465.0
3. Product ID: B007HSKSMI, Sum of Ratings: 1458.0
4. Product ID: B00KXVBB3Q, Sum of Ratings: 1171.0
5. Product ID: B0151K2AB0, Sum of Ratings: 1114.0
6. Product ID: B00JPBFC8U, Sum of Ratings: 842.0
7. Product ID: B005SN3INA, Sum of Ratings: 803.0
8. Product ID: B00R0RBPC0, Sum of Ratings: 564.0
9. Product ID: B0006U3ACY, Sum of Ratings: 517.0
10. Product ID: B00KMRVGF0, Sum of Ratings: 496.0
```